

# Indexing


- ↳ Need (why)
- ↳ Implementation (How)

- ↳ How file system indexing is different from DBMS indexing

Many queries reference only a small proportion of the records in a file. For example, a query like "Find all instructors in the Physics department" or "Find the total number of credits earned by the student with ID 22201" references only a fraction of the student records. It is inefficient for the system to read every tuple in the *instructor* relation to check if the *dept\_name* value is "Physics". Likewise, it is inefficient to read the entire *student* relation just to find the one tuple for the ID "32556". Ideally, the system should be able to locate these records directly. To allow these forms of access, we design additional structures that we associate with files.

\* Index is also a file but it doesn't contain actual data but pointer to data.





↓  
Index file  
having ID as  
Index

\* Both actual file of records and index file associated with it are present in HDD.

\* When DBMS process is created, it loads index file into memory for faster look-ups.

\* Important part is what data structure can be used to

make search on this Index  
file faster

↳ Tree  $\rightarrow \log(n)$  complexity

\* Size of Index file  $\ll \ll$  Size  
of actual records

||  
This will take  
less HDD blocks than this

\* So, accessing time/ searching time  
is significantly decreased.

$\Rightarrow$ 

Key	HDD block pointer
-----	-------------------

↳ Each row in Index file

\* One table could have multiple indexes/index file depending on type of SQL query

#### 1. Purpose of Indexes:

- Speed up SELECT queries and WHERE clauses.
- Improve the performance of sorting and grouping operations, as well as JOINS.

#### 2. How Indexes Work:

- Indexes are built on columns of a table. When a query searches for values in these columns, the database engine uses the index instead of scanning row by row.
- Indexes often use structures like B-Trees or Hashes. B-Tree indexes are most common because they allow ordered traversal, which is useful for range-based queries.

#### 4. Index Costs:

- Storage Overhead: Indexes consume additional disk space.
- Slower Writes: Inserts, updates, and deletes can slow down since indexes also need to be updated.
- **Maintenance**: Regular updates or high data modification rates can require index maintenance.

#### 5. Index Best Practices:

- Index columns frequently used in WHERE clauses, JOINS, and ORDER BY clauses.
- Avoid over-indexing; too many indexes can slow down write operations and increase storage costs.
- Consider composite indexes for columns often queried together.
- Periodically monitor and adjust indexes based on query performance analysis.