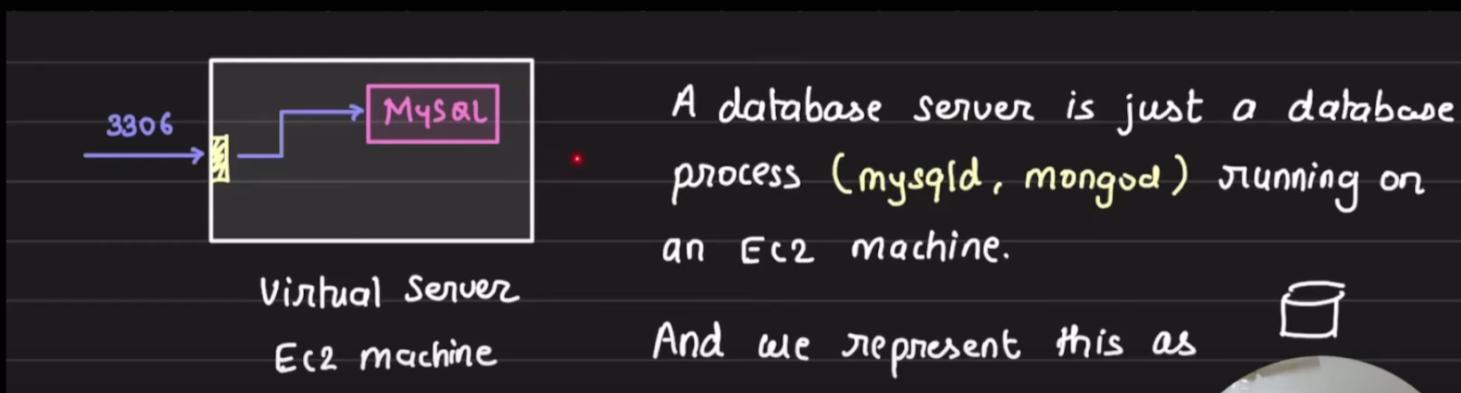


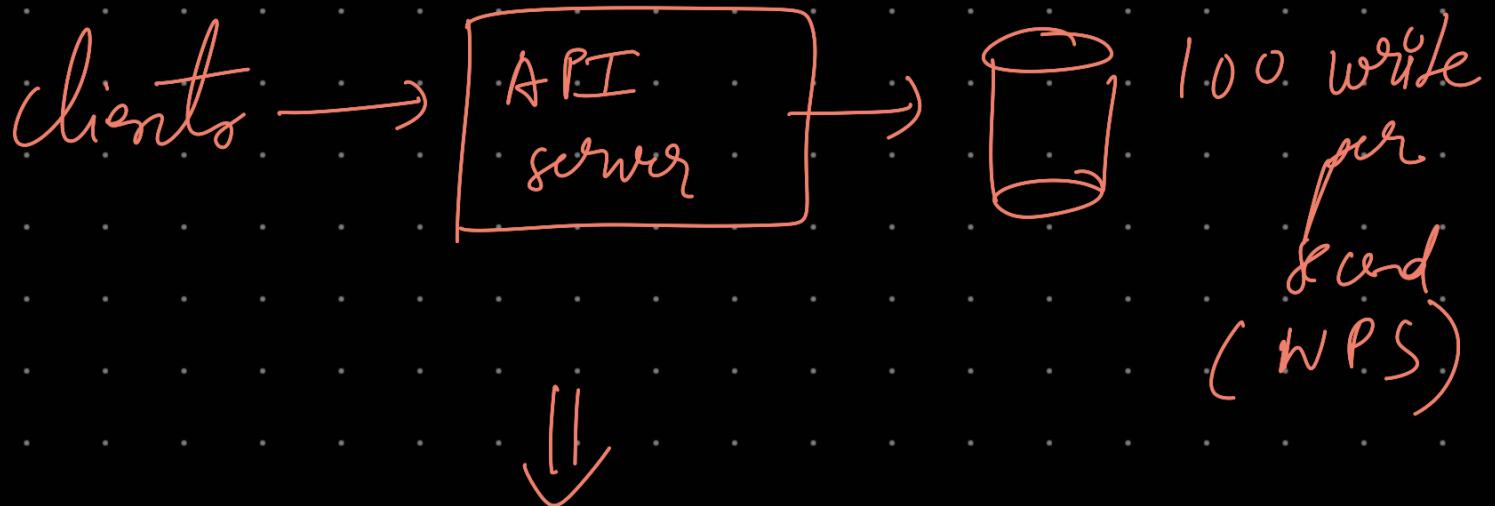
Database

→ It's just a process running
on some port in a physical
server machine



Let say I wrote a S/W with
a Database and now I put it
on production using some cloud provider





We started getting more users



Database metrics drops down

* CPU ↑ RAM ↑

* Query time increased



Vertically scale

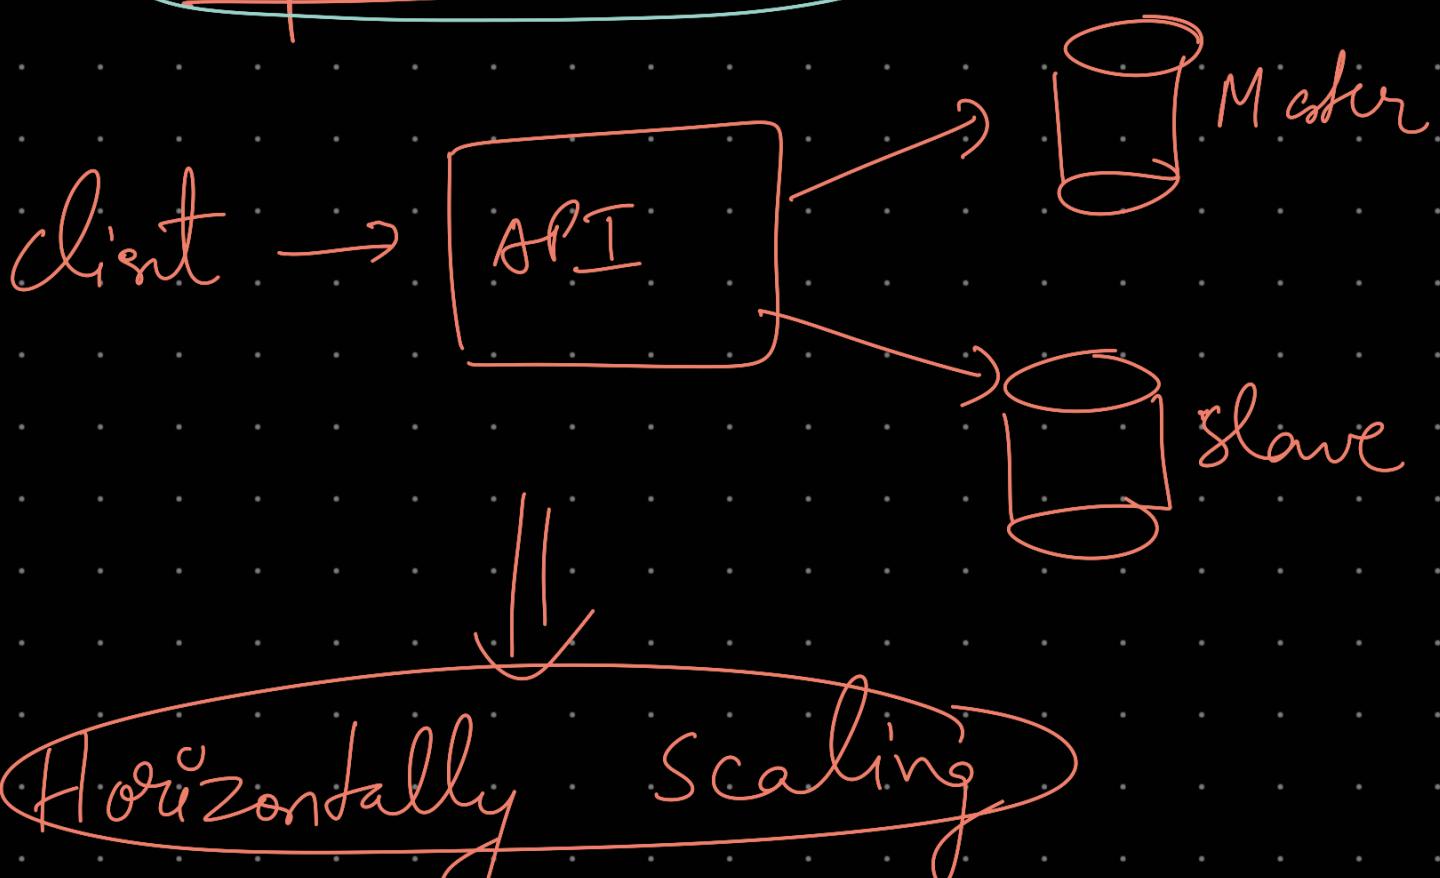
(Add more RAM, CPU, Disk)



Single Machine reaches Max limit



Replicate Data



Horizontally Scaling

We also **split** data
→ partitioning

Data level → partitioning

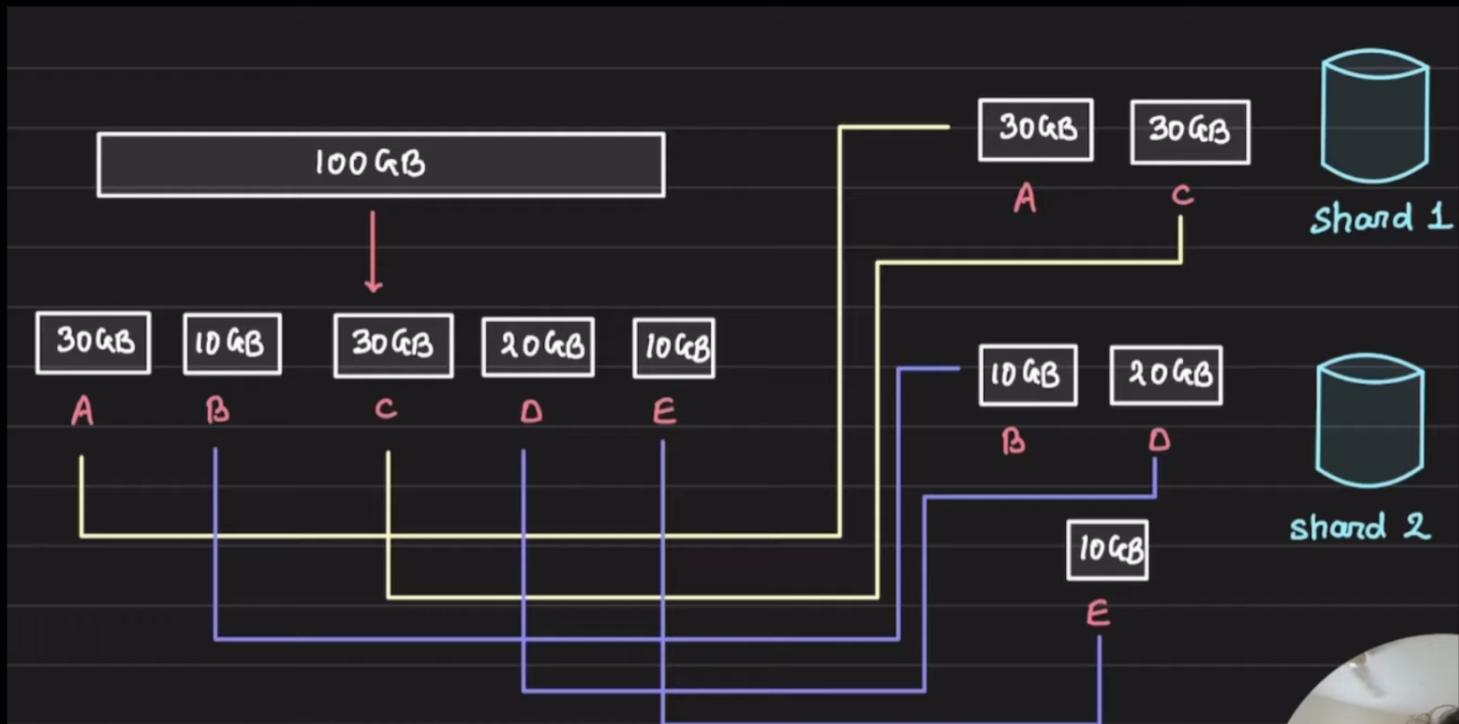
Machine level
or
Database
Level → sharding

* Each database is a shard
and we partitioned the data



You partitioned the 100 GB of total data into 5 mutually exclusive partitions.

Each of these partitions can either live on one database server or a couple of them can share one server.



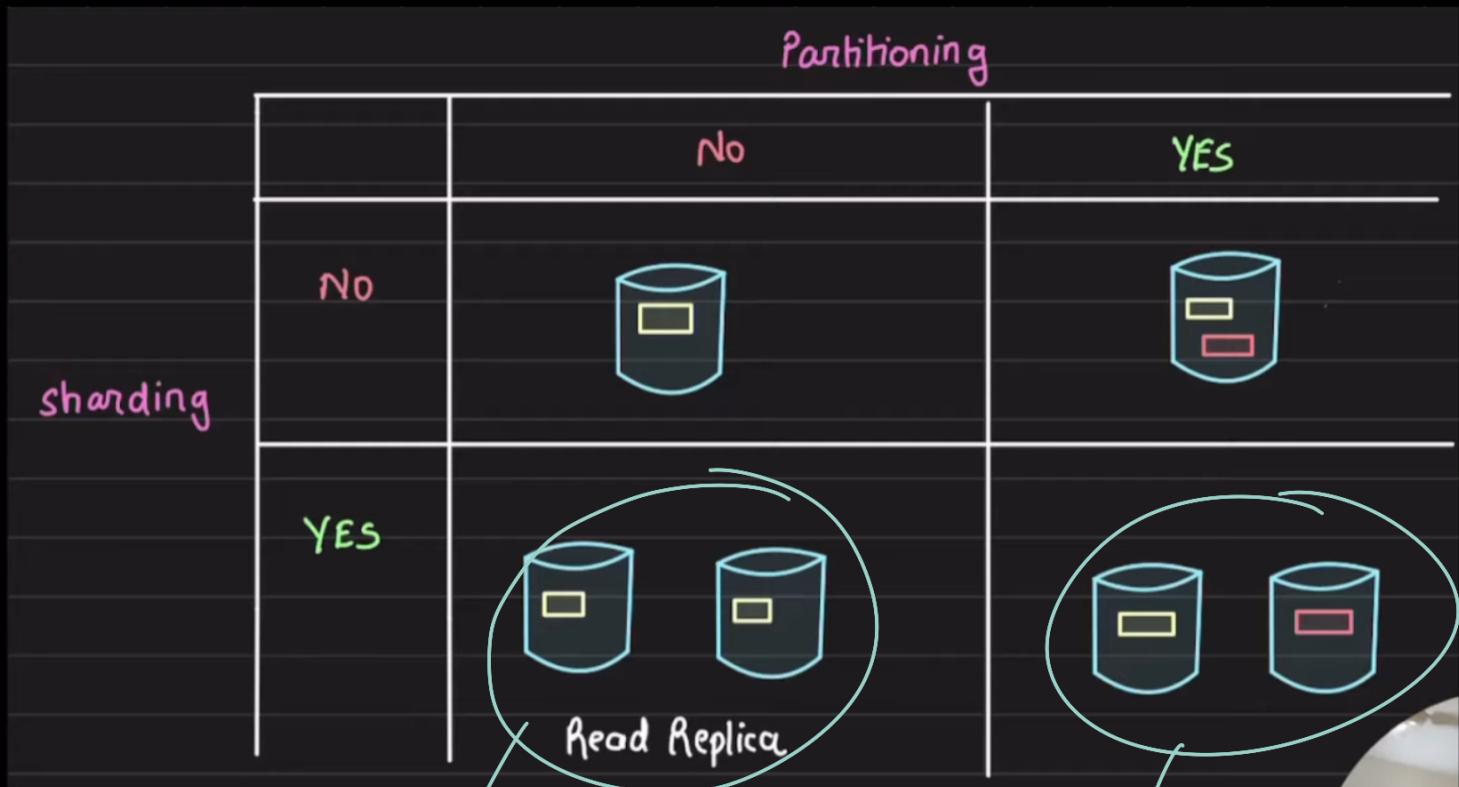
Partitioning

Horizontal

Row placement
splitting

Vertical

Column splitting



High Reads

High reads + writer

Advantages of sharding

- Handle large reads and writes
- Increase overall storage capacity
- Higher availability

Disadvantages of sharding

- operationally complex
- cross-shard queries expensive

* In case of JOINS across shards/Machine, networking is required → Latency increased

~~Steps~~

- ① Check slow reads / write
- ② Analyze slowest queries
- ③ Create index
- ④ Introduce Cache
- ⑤ Replicate DB
- ⑥ Scale vertically

Exhausted all option, then

↓

Horizontal scaling + sharding

Compromise on ACID